

压测监控及pod调优

一、压测流程

1、确认压测场景

参会的人员：架构师、后端、测试

确认事项：

- 1、压测场景
- 2、压测目标
- 3、中间件配置

2、压测方案评审

参会人员：架构师、后端、测试

3、发邮件通知pm及运维购买相关资源

4、开发人员搭建压测环境及数据库迁移

5、开发人员提供压测数据

6、运维人员给出监控页面

7、开始压测

二、概念介绍

1、压测的目的

压力测试是对系统不断施加压力，来获得系统最大服务能力的测试。一般而言，只有在系统基础功能测试验证完成、系统趋于稳定的情况下，才会进行压力测试。

当负载逐渐增加时，观察系统各项性能指标的变化情况是否有异常。

2、压测指标

压测的指标通常有tps、并发数、响应时间、错误率等

1) tps每秒事务处理量(TransactionPerSecond)

每秒处理的消息数。

2) qps每秒处理事务数

包括了用户请求服务器、服务器自己的内部处理、服务器返回给用户。这三个过程，每秒能够完成N个这三个过程，Tps也就是N。

3) qps与tps关系

Qps基本类似于Tps，但是不同的是，对于一个页面的一次访问，形成一个Tps；但一次页面请求，可能产生多次对服务器的请求，服务器对这些请求，就可计入“Qps”之中。

例如，访问一个 Index 页面会请求服务器 3 次，包括一次 html，一次 css，一次 js，那么访问这一个页面就会产生一个“T”，产生三个“Q”。

3) 并发数

系统同时能处理的请求数量

这里有两个概念介绍（用户端并发与服务端并发）


举个例子：

一款商品，晚上八点订阅用户有10w，此时用户端并发就为10w。由于不同的用户使用的硬件不同以及点击下单的时间差，导致并不是同一时间触发请求（假设这里损失1w的用户），接下来会有9w的用户同时请求了服务端，又由于每位用户的地理位置不同，假设服务器位于北京，靠近北京的用户当然比远在海南的用户先到服务端，这里再损失3w的用户。此时服务端并发就为6w了。所以服务端并发<用户端并发

三、压测关注点

1、cpu监控


微服务cpu监控



服务pod cpu监控

公共-测试环境 / TKE-公共-测试环境资源监控 ☆

prod-gw-wechats-prod-common-queue-86966df57f-g67jk	-	0.90
prod-micro-goods-v2-prod-common-queue-6f6d7f5856-zbb9m	-	0.02
prod-micro-activity-v2-prod-common-queue-557b5d6987-dnktg	-	0.03
prod-micro-order-v2-prod-common-queue-	-	0.03




容器内存使用量(可关联节点)

2023-02-16 11:15:15

prod-gw-shop-v2-prod-common-http-68645d768-jg9w8:	71.48%
prod-gw-shop-v2-prod-common-http-68645d768-2wqpr:	70.30%
prod-gw-shop-v2-prod-common-http-68645d768-rkj6n:	69.85%
prod-gw-shop-v2-prod-common-http-68645d768-vzs2d:	69.78%
prod-gw-shop-v2-prod-common-http-68645d768-n97hs:	69.64%
prod-gw-shop-v2-prod-common-http-68645d768-rsgdr:	69.60%
prod-gw-shop-v2-prod-common-http-68645d768-zktp:	69.53%
prod-gw-shop-v2-prod-common-http-68645d768-cwj6k:	69.03%
prod-gw-shop-v2-prod-common-http-68645d768-9v6px:	68.78%
prod-gw-shop-v2-prod-common-http-68645d768-mfzkb:	68.22%
prod-gw-shop-v2-prod-common-http-68645d768-g9r9b:	68.21%
prod-gw-shop-v2-prod-common-http-68645d768-6d4qp:	68.13%
prod-gw-shop-v2-prod-common-http-68645d768-7htcj:	60.46%
prod-gw-shop-v2-prod-common-http-68645d768-scsq5:	60.37%
prod-gw-shop-v2-prod-common-http-68645d768-kcmns:	59.98%
prod-gw-shop-v2-prod-common-http-68645d768-9p9kp:	59.33%
prod-gw-shop-v2-prod-common-http-68645d768-zg89s:	59.27%
prod-gw-shop-v2-prod-common-http-68645d768-m8wwq:	59.22%
prod-micro-goods-v2-prod-common-http-7dc9986d86-lmbjb:	58.54%
prod-micro-goods-v2-prod-common-http-7dc9986d86-4kn7q:	58.51%
prod-micro-goods-v2-prod-common-http-7dc9986d86-kfjd5:	58.29%
prod-micro-goods-v2-prod-common-http-7dc9986d86-zbv82:	58.04%
prod-gw-shop-v2-prod-common-http-68645d768-qcqlm:	39.89%
prod-gw-shop-v2-prod-common-http-68645d768-gcltq:	38.71%
prod-gw-shop-v2-prod-common-http-68645d768-qz7zp:	38.50%
prod-gw-shop-v2-prod-common-http-68645d768-68twb:	38.44%
prod-micro-goods-v2-prod-common-http-7dc9986d86-hqjks:	34.18%
prod-gw-shop-v2-prod-common-http-68645d768-s4scj:	33.14%
prod-gw-shop-v2-prod-common-http-68645d768-q65j9:	32.93%
prod-micro-activity-v2-prod-common-http-75bdb746cc-d4rb2:	30.70%
prod-micro-activity-v2-prod-common-http-75bdb746cc-w5p2h:	30.46%
prod-gw-shop-v2-prod-common-http-68645d768-v9zic:	27.54%

cvm单机 cpu监控

CPU 利用率

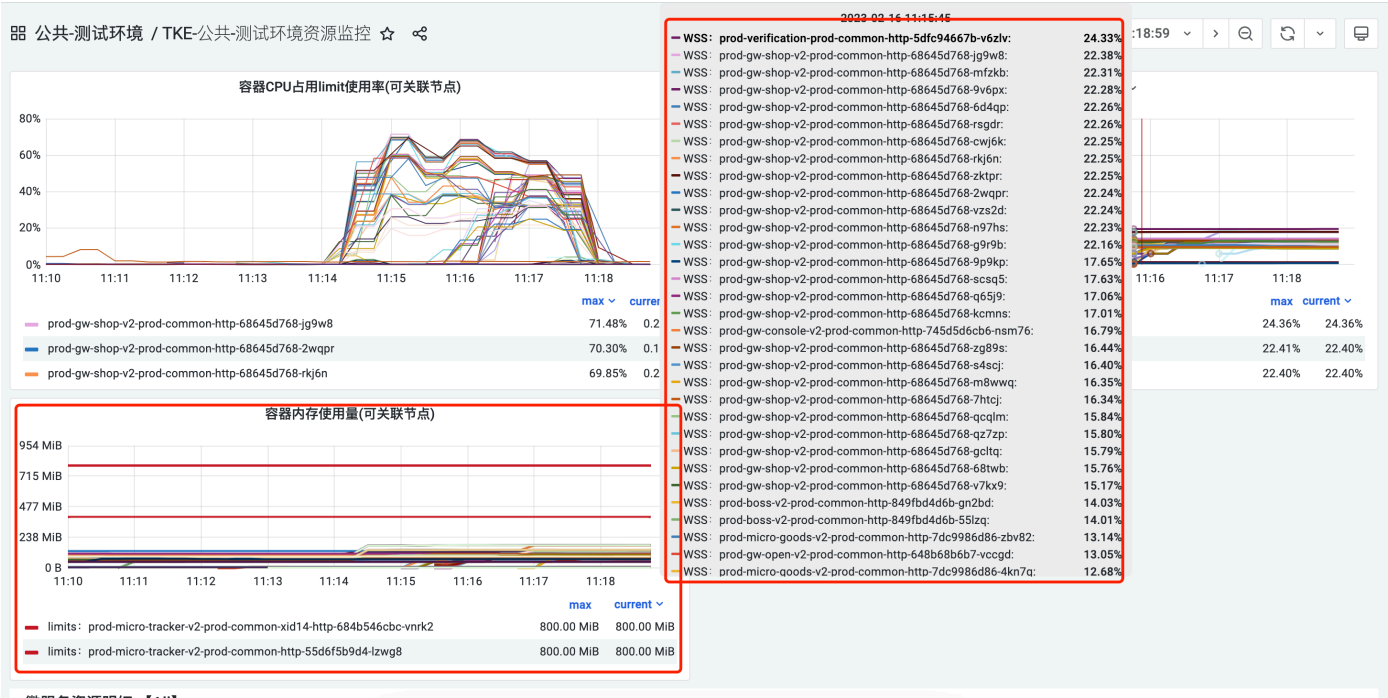


2、内存监控

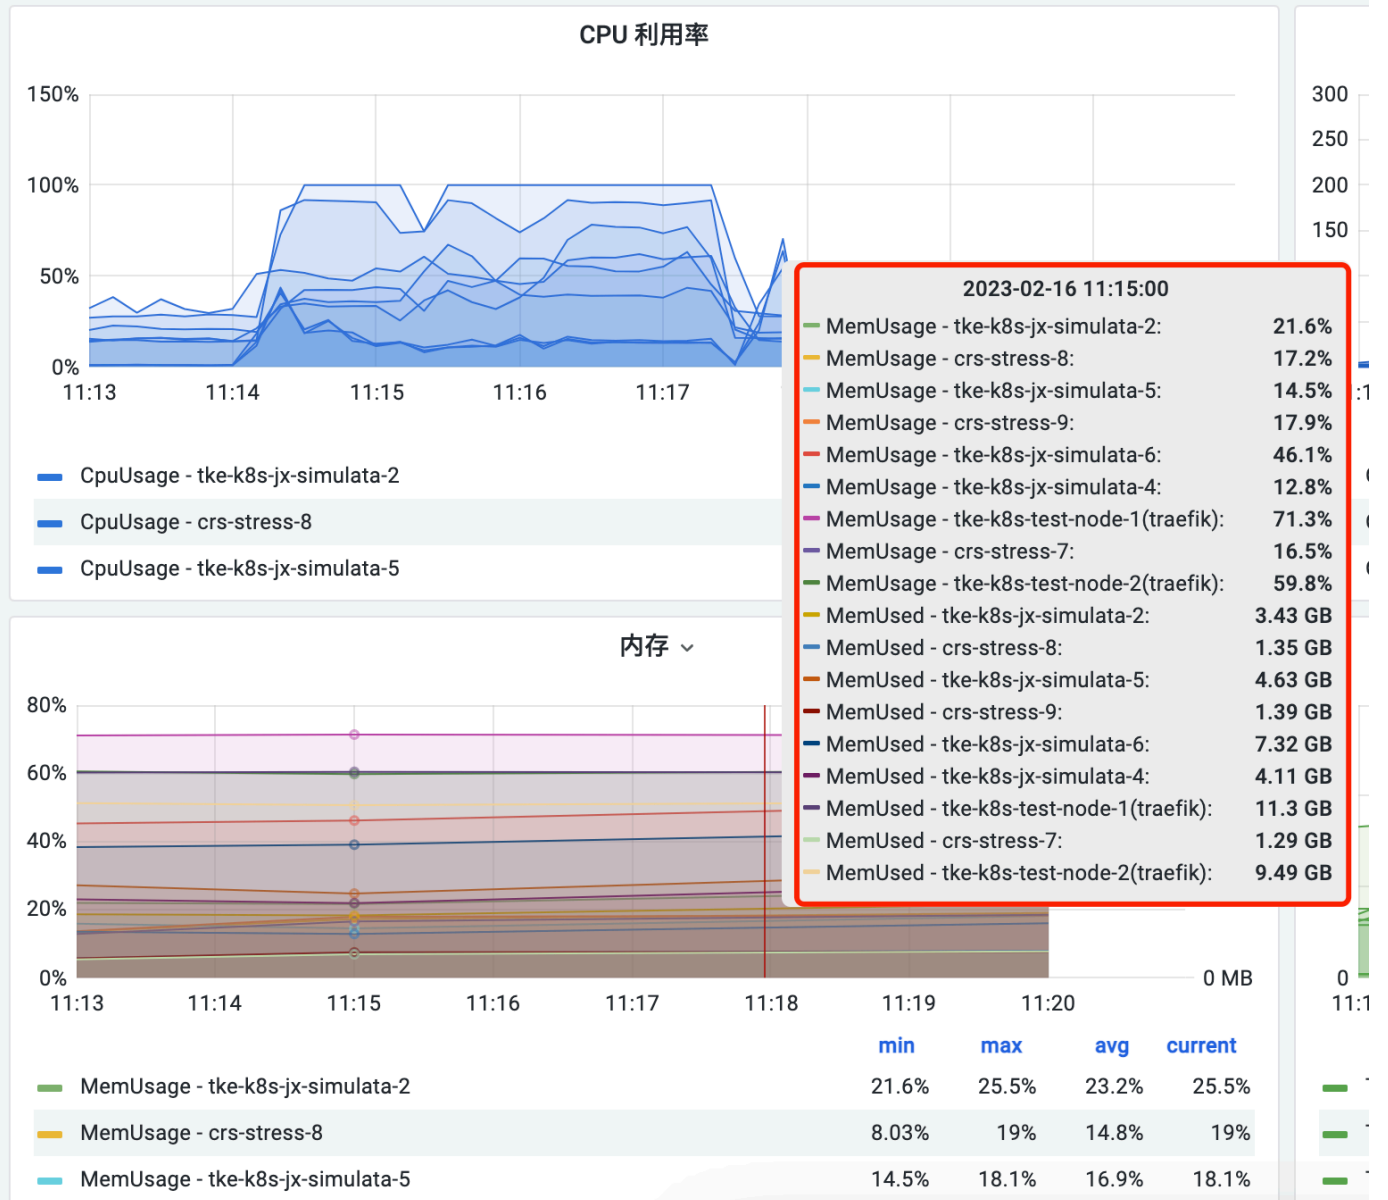
微服务内存监控



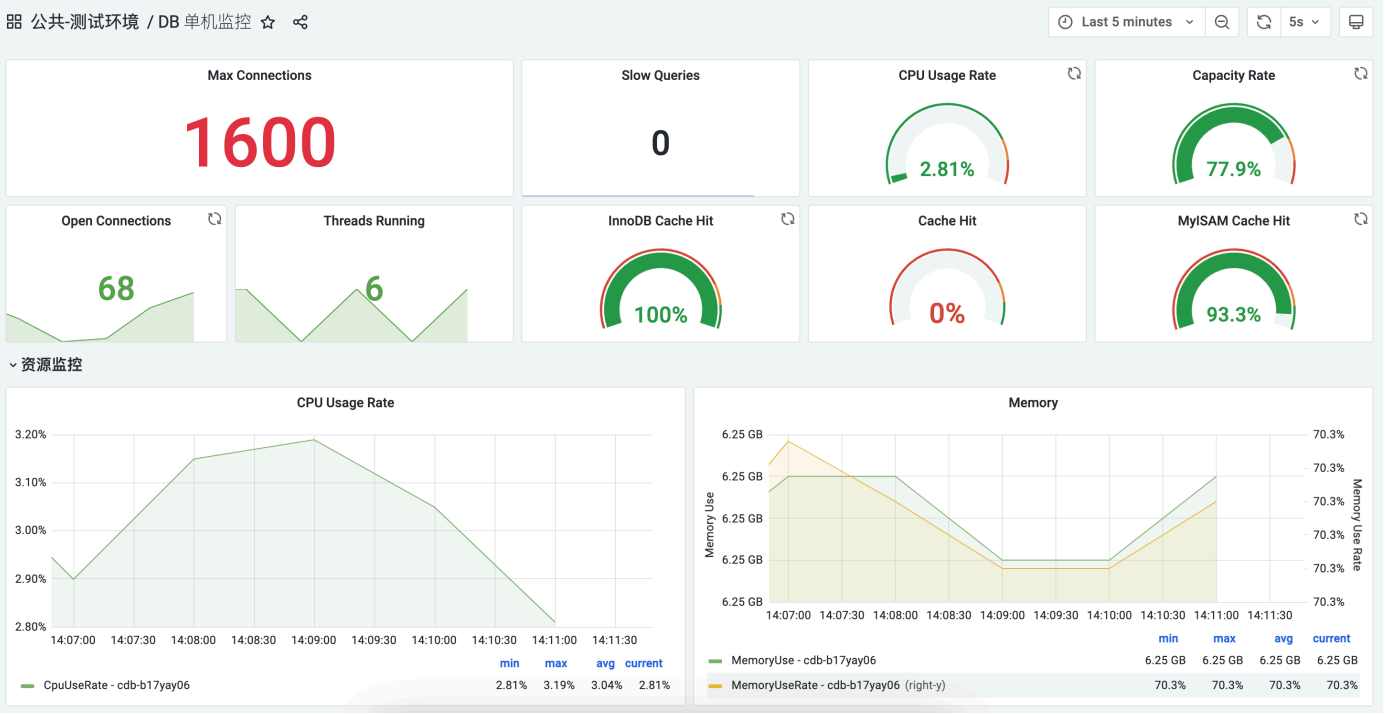
服务pod 内存监控



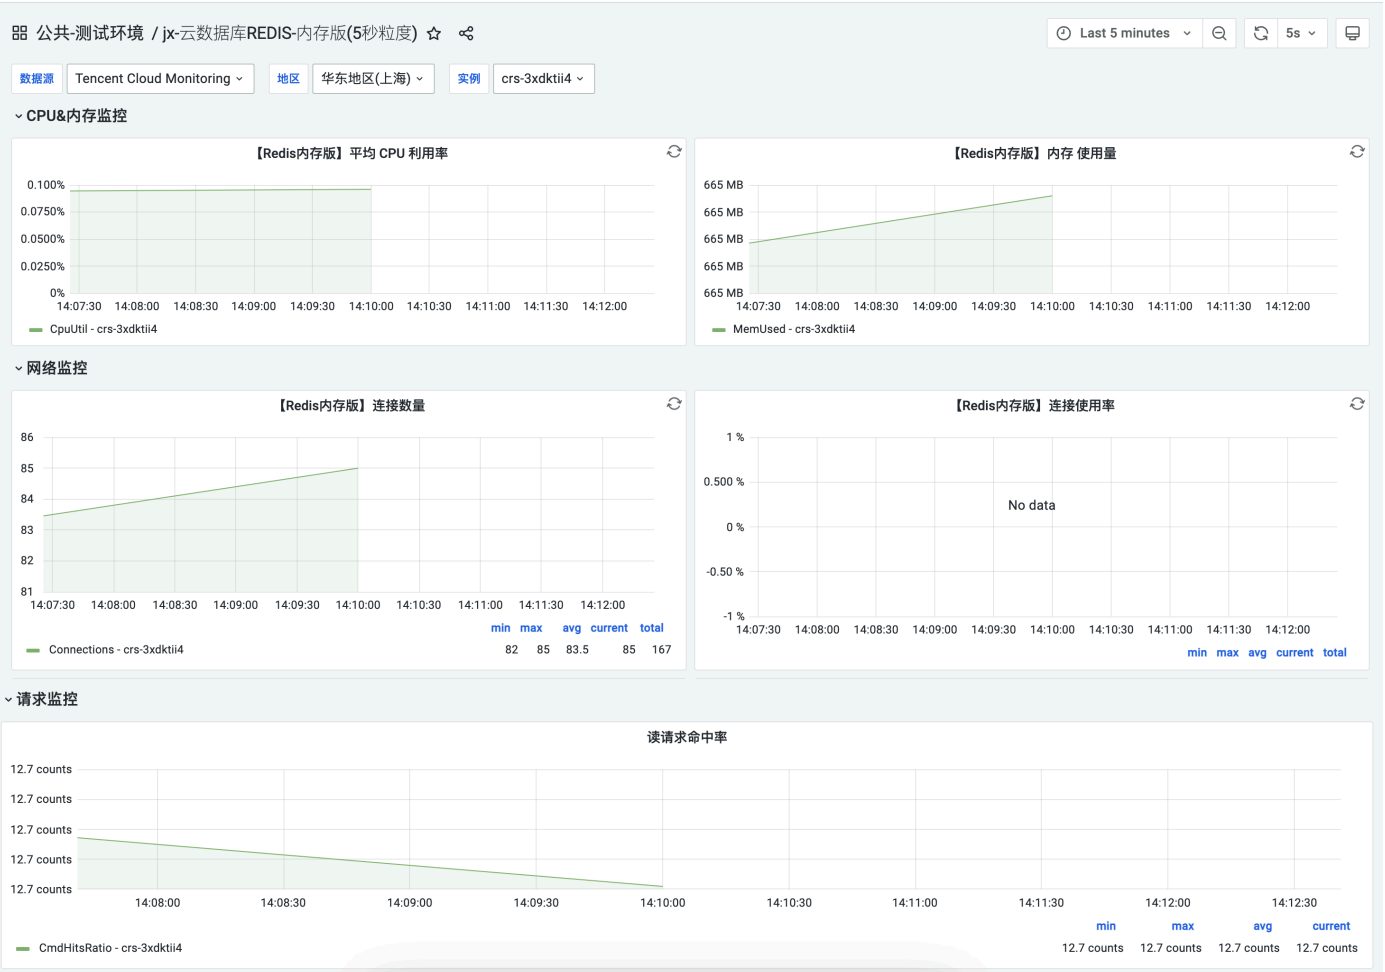
cvm单机内存监控



3、观察DB使用情况



4、观察Redis使用情况



四、监控流程

1、制作表格

服务	pod		cpu		memory		cpu (总数)			memory (总数)	
	request	limit	request	limit	request	limit	request(pod)*request(cpu)	limit(pod)*request(cpu)	request(pod)*limit(cpu)	request	limit
gw-console-v2	1	1	0.5	0.8	500	800	0.5	0.5	0.8	500	500
gw-shop-v2	15	20	0.5	0.8	500	800	7.5	10	12	7500	10000
boss-v2	2	2	0.5	0.8	500	800	1	1	1.6	1000	1000
gw-open-v2	1	1	0.5	0.8	500	800	0.5	0.5	0.8	500	500
micro-goods-v2	5	10	0.5	0.8	500	800	2.5	5	4	2500	5000
micro-order-v2	1	1	0.5	0.8	500	800	0.5	0.5	0.8	500	500
micro-activity-v2	5	10	0.5	0.8	500	800	2.5	5	4	2500	5000
go-micro-user	1	1	0.3	0.5	500	800	0.3	0.3	0.5	500	500
gw-wechats	2	2	0.3	0.5	500	800	0.6	0.6	1	1000	1000
micro-download-v2	1	1	0.3	0.5	500	800	0.3	0.3	0.5	500	500
micro-tracker	1	1	0.3	0.5	500	800	0.3	0.3	0.5	500	500
总数			32		64000		16.5	24	26.5	17500	25000

表格方便计算每次调整pod数、cpu数、内存数。其中H列与J列需小于总CPU核数（32），相对应的内存计算应小于总内存数（64G）。

如下图，每个cvm主机不光运行了业务pod外，还运行了基础服务，这些也是会占用cpu及内存，所以需要预留cpu与内存给基础服务。

监测pod所属cvm

Namespace	Name	CPU Requests	CPU Limits	Memory Requests	Memory Limits	AGE
argocd	argocd-reis-ha-haproxy-5bc5dbc689-6jj24	0 (0%)	0 (0%)	0 (0%)	0 (0%)	35d
argocd	argocd-repo-server-765b6445cd-zrkdK	0 (0%)	0 (0%)	0 (0%)	0 (0%)	35d
argocd	argocd-server-6f8d69c7dc-4w6s6	0 (0%)	0 (0%)	0 (0%)	0 (0%)	35d
hivesec	hiveagent-qzqxl	500m (6%)	1 (12%)	500Mi (1%)	1000Mi (3%)	2d20h
infra	rocketmq-dashboard-774c79ddb9-d5b7t	0 (0%)	0 (0%)	0 (0%)	0 (0%)	19d
kube-system	csi-coslauncher-xrsh6	250m (3%)	8 (102%)	250Mi (0%)	8Gi (28%)	36d
kube-system	csi-cosplugin-7kb2k	0 (0%)	0 (0%)	0 (0%)	0 (0%)	36d
kube-system	csi-nodeplugin-cfsplugin-7hkhP	0 (0%)	0 (0%)	0 (0%)	0 (0%)	36d
kube-system	ip-masq-agent-nq2l5	0 (0%)	0 (0%)	0 (0%)	0 (0%)	36d
kube-system	kube-proxy-nbjkl	0 (0%)	0 (0%)	0 (0%)	0 (0%)	36d
kube-system	log-pilot-lwv9r	0 (0%)	0 (0%)	0 (0%)	0 (0%)	8d
kube-system	node-problem-detector-fdrfd	100m (1%)	500m (6%)	200Mi (0%)	200Mi (0%)	36d
kube-system	serf-agent-9f4v9	0 (0%)	0 (0%)	0 (0%)	0 (0%)	36d
kube-system	tke-bridge-agent-n9tg4	0 (0%)	0 (0%)	0 (0%)	0 (0%)	36d
kube-system	tke-cni-agent-8ttgs	0 (0%)	0 (0%)	0 (0%)	0 (0%)	36d
kube-system	tke-eni-agent-mxjsb	10m (0%)	0 (0%)	0 (0%)	0 (0%)	36d
kube-system	tke-monitor-agent-fxdph	0 (0%)	100m (1%)	0 (0%)	100Mi (0%)	36d
kubesphere-monitoring-system	node-exporter-lf5kp	112m (1%)	2 (25%)	200Mi (0%)	600Mi (2%)	36d
shopping-mall-baseline-prod	prod-boss-v2-prod-common-http-849fbd4d6b-gn2bd	300m (3%)	500m (6%)	500Mi (1%)	800Mi (2%)	10m
shopping-mall-baseline-prod	prod-go-micro-user-prod-common-http-5755744b55-s4z8x	100m (1%)	100m (1%)	500Mi (1%)	800Mi (2%)	9m32s
shopping-mall-baseline-prod	prod-gw-open-v2-prod-common-http-648b68b6b7-vccgd	500m (6%)	800m (10%)	500Mi (1%)	800Mi (2%)	20h
shopping-mall-baseline-prod	prod-gw-open-v2-prod-common-queue-5b46c699bf-h6lhm	0 (0%)	0 (0%)	0 (0%)	0 (0%)	20h
shopping-mall-baseline-prod	prod-gw-shop-v2-prod-common-http-68645d768-47svc	300m (3%)	800m (10%)	500Mi (1%)	800Mi (2%)	6m15s
shopping-mall-baseline-prod	prod-gw-shop-v2-prod-common-http-68645d768-4b6rc	300m (3%)	800m (10%)	500Mi (1%)	800Mi (2%)	6m15s
shopping-mall-baseline-prod	prod-gw-shop-v2-prod-common-http-68645d768-7rcwf	300m (3%)	800m (10%)	500Mi (1%)	800Mi (2%)	6m15s
shopping-mall-baseline-prod	prod-gw-shop-v2-prod-common-http-68645d768-fnz92	300m (3%)	800m (10%)	500Mi (1%)	800Mi (2%)	6m15s
shopping-mall-baseline-prod	prod-gw-shop-v2-prod-common-http-68645d768-hhw2f	300m (3%)	800m (10%)	500Mi (1%)	800Mi (2%)	6m15s
shopping-mall-baseline-prod	prod-gw-shop-v2-prod-common-http-68645d768-q48n2	300m (3%)	800m (10%)	500Mi (1%)	800Mi (2%)	6m15s
shopping-mall-baseline-prod	prod-gw-shop-v2-prod-common-http-68645d768-qcqlm	300m (3%)	800m (10%)	500Mi (1%)	800Mi (2%)	15m
shopping-mall-baseline-prod	prod-gw-shop-v2-prod-common-http-68645d768-v9zjc	300m (3%)	800m (10%)	500Mi (1%)	800Mi (2%)	6m15s
shopping-mall-baseline-prod	prod-gw-wechats-prod-common-http-75b7c4bf4f-psjrc	300m (3%)	300m (3%)	500Mi (1%)	800Mi (2%)	9m50s
shopping-mall-baseline-prod	prod-micro-activity-v2-prod-common-http-75bdb746cc-4cpqx	300m (3%)	800m (10%)	500Mi (1%)	800Mi (2%)	10m
shopping-mall-baseline-prod	prod-micro-activity-v2-prod-common-http-75bdb746cc-l62pm	300m (3%)	800m (10%)	500Mi (1%)	800Mi (2%)	4m56s
shopping-mall-baseline-prod	prod-micro-activity-v2-prod-common-http-75bdb746cc-nw4mm	300m (3%)	800m (10%)	500Mi (1%)	800Mi (2%)	11m
shopping-mall-baseline-prod	prod-micro-goods-v2-prod-common-http-7dc9986d86-65kxm	300m (3%)	800m (10%)	500Mi (1%)	800Mi (2%)	6m44s
shopping-mall-baseline-prod	prod-micro-goods-v2-prod-common-http-7dc9986d86-hqjks	300m (3%)	800m (10%)	500Mi (1%)	800Mi (2%)	12m
shopping-mall-baseline-prod	prod-micro-tracker-v2-prod-common-xid14-http-684b546cbc-vnrk2	100m (1%)	100m (1%)	500Mi (1%)	800Mi (2%)	9m7s
Allocated resources:						
(Total limits may be over 100 percent, i.e., overcommitted.)						
Resource	Requests	Limits				
cpu	6172m (79%)	23800m (305%)				
memory	10150Mi (35%)	24492Mi (86%)				
ephemeral-storage	0 (0%)	0 (0%)				
hugepages-1Gi	0 (0%)	0 (0%)				
hugepages-2Mi	0 (0%)	0 (0%)				
tke.cloud.tencent.com/eip	0	0				
tke.cloud.tencent.com/eni-ip	0	0				
Events:	<none>					

2、估算所压场景所涉及的服务

例如：

压测确认订单场景，梳理代码可得出涵盖的服务范围有（shop、activity、goods、order、boss、user）。

其中shop为网关入口，order主要负责逻辑运算、activity、goods主负责查询业务、user与boss则为辅助服务。此时可以得出order的cpu相对于shop、activity、goods高一点，shop的pod数相对于order、activity、goods高点，user与boss因有缓存，给出对应的两个则行。

最终第一次压测所得出的结论：shop给出8-10个、order与goods与activity给出5个pod、user与boss给出1个。

3、根据第一次压测结果调整修改pod

观察cpu、内存、DB以及Redis的数据（因此次压测中出现的瓶颈问题都为cpu，故在这里只分析根据cpu的调整）

4、常见问题分析

- 1) 并发小、错误率大。首先排查脚本自身原因与业务层面的原因。
- 2) 响应时常长、TPS小。检查压力机资源消耗，首先确保不是压力机的性能原因导致相应时常长。
- 3) 并发增加、cpu压不上去。排查数据库、php-fpm、nginx等原因。
- 4) TPS低、响应时常长。确认业务流向，例如：jmeter->nginx->php-fpm->php->缓存->mysql->返回。